




Figure 1: Max-cut: certified approximation ratio of random greedy (RG, mean over trials) with respect to each upper bound, as a function of the cardinality k . $f(S)$ = weight of edges crossing $(S, V \setminus S)$; $G(n, p)$ with integer weights in $\{1, \dots, 9\}$. RG/OPT uses the exact optimum (brute force) where it was computable. The total-weight bound is shown only where it is within $1.5\times$ the dual bound. Curves are means over graph/sample seeds.

Random greedy: certified approximation ratio vs. k

Each plot shows RG/B for the dual upper bound, the top- k singleton bound and (where informative) the total-weight bound B , against the $1/e$ guarantee of random greedy.




Figure 2: Directed cut: certified approximation ratio of random greedy (RG, mean over trials) with respect to each upper bound, as a function of the cardinality k . $f(S)$ = weight of edges leaving S ; synthetic $G(n, p)$ with random orientations and SNAP graphs. RG/OPT uses the exact optimum (brute force) where it was computable. The total-weight bound is shown only where it is within $1.5\times$ the dual bound. On the largest instances the dual bound is minimized over every 10th greedy prefix only. Curves are means over graph/sample seeds.




Figure 3: Revenue maximization: certified approximation ratio of random greedy (RG, mean over trials) with respect to each upper bound, as a function of the cardinality k . $f(S) = \sum_{i \notin S} \sqrt{\sum_{j \in S} w_{ij}}$ with $w_{ij} \sim U(0, 1)$. RG/OPT uses the exact optimum (brute force) where it was computable. The total-weight bound is shown only where it is within $1.5\times$ the dual bound. On the largest instances the dual bound is minimized over every 5th or 10th greedy prefix only. Curves are means over graph/sample seeds.




Figure 4: Diverse recommendation: certified approximation ratio of random greedy (RG, mean over trials) with respect to each upper bound, as a function of the cardinality k . $f(S) = \sum_{i \in S} \sum_{j \in V} s_{ij} - \lambda \sum_{i, j \in S} s_{ij}$, cosine similarities of MovieLens 1M rating vectors. RG/OPT uses the exact optimum (brute force) where it was computable. Curves are means over graph/sample seeds.




Figure 5: Log-determinant: certified approximation ratio of random greedy (RG, mean over trials) with respect to each upper bound, as a function of the cardinality k . $f(S) = \log \det(L_S)$, $L = \alpha(K + 10^{-3}I)$ with an RBF kernel K on the features. RG/OPT uses the exact optimum (brute force) where it was computable. On the largest instances the dual bound is minimized over every 5th greedy prefix only. Curves are means over graph/sample seeds.




Figure 6: Gaussian mutual information: certified approximation ratio of random greedy (RG, mean over trials) with respect to each upper bound, as a function of the cardinality k . $f(S) = I(X_S; X_{V \setminus S})$ for a Gaussian process on random points and for the Intel Berkeley lab temperature sensors. RG/OPT uses the exact optimum (brute force) where it was computable. Curves are means over graph/sample seeds.




Figure 7: Hypergraph cut: certified approximation ratio of random greedy (RG, mean over trials) with respect to each upper bound, as a function of the cardinality k . $f(S)$ = weight of hyperedges split by S ; random hypergraphs and the MAG-10 co-authorship hypergraph. RG/OPT uses the exact optimum (brute force) where it was computable. The total-weight bound is shown only where it is within $1.5\times$ the dual bound. Curves are means over graph/sample seeds.